DETECTING CANCER METASTASES ON GIGAPIXEL PATHOLOGY IMAGE

XINLEI CAO UNI: xc2420

GOAL OF PROJECT

- Each year, the treatment decisions for more than 230,000 breast cancer patients in the U.S. hinge on whether the cancer has metastasized away from the breast. Metastasis detection is currently performed by pathologists reviewing large expanses of biological tissues. This process is labor intensive and error-prone. (copied from Paper: Detecting Cancer Metastases on Gigapixel Pathology Images)
- So we want to build a model to automatically detect and localize tumors.

INTRODUCTION TO MY METHOD

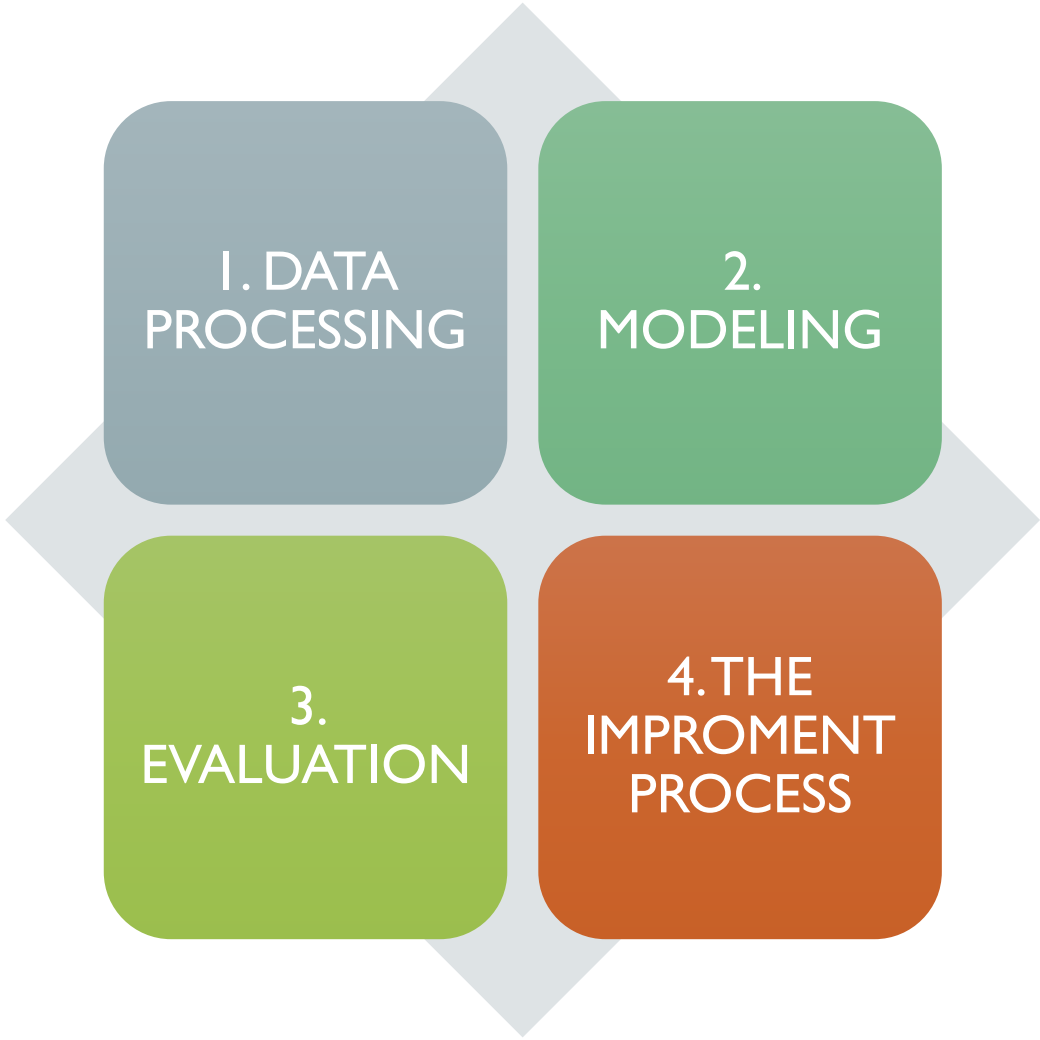

DATA PROCESSING

- **1. USING SLIDE WINSOW TO GENERATE IMAGE PATCH**
- Slide Level = 3
- Patch Size = 150, stride = 80, Center Size = 100

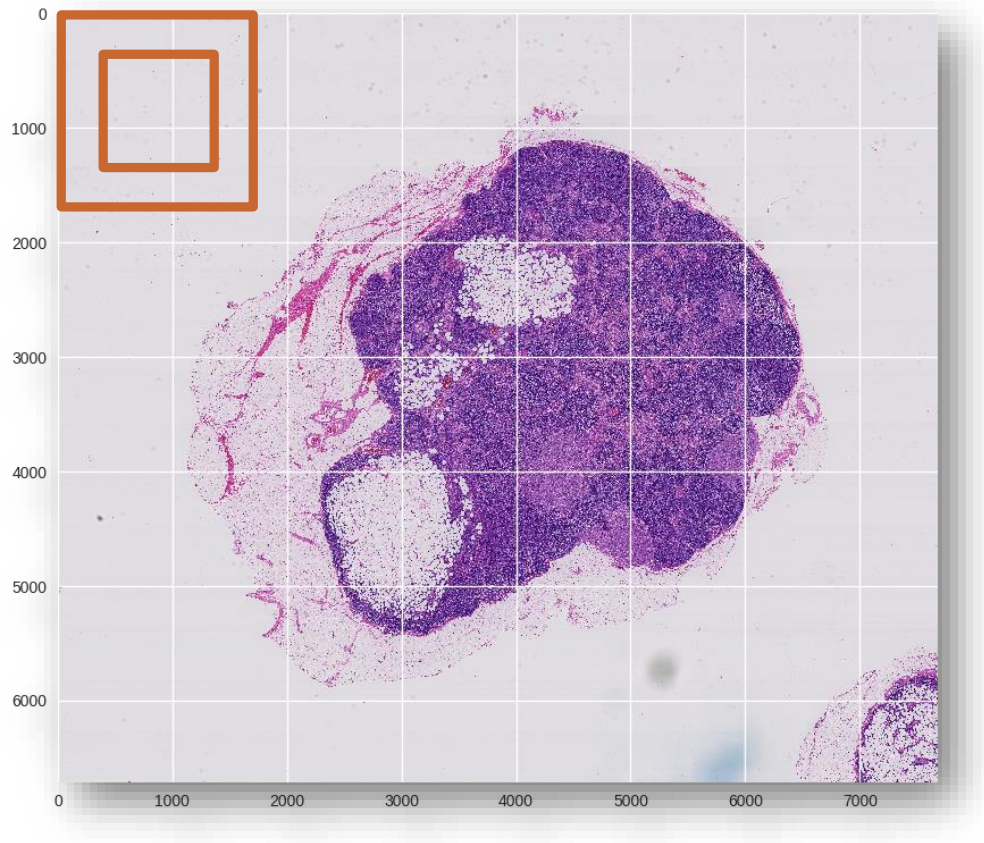

DATA PROCESSING

- **2. DATA AUGMENTATION**
- Flip and Rotation
- Only applied on the tumor patch
- 209 tumor images → 1672 tumor images

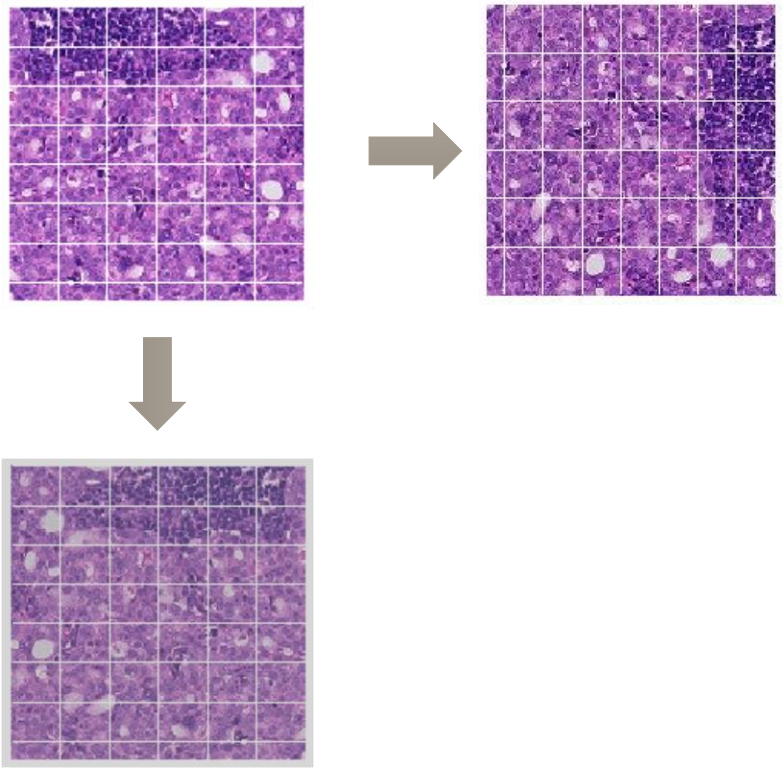

DATA PROCESSING

- **3. IGNORE NON-TISSUE AREAS**
- Ignore the patches whose tissue proportion is less than 0.2

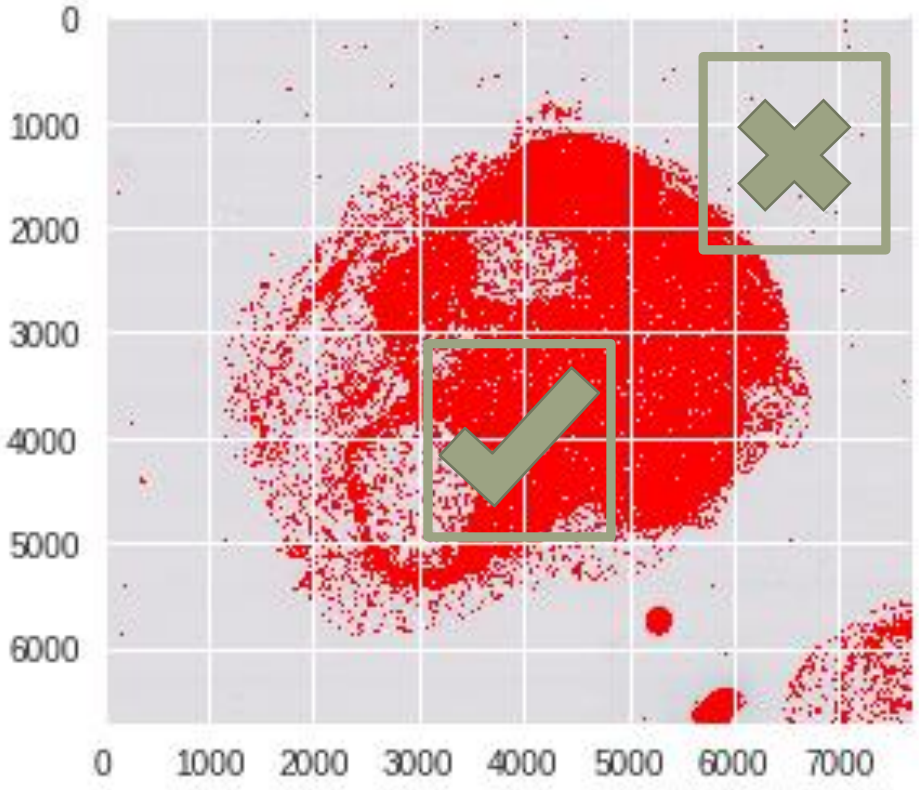

DATA PROCESSING

- **4. HANDEL UNBALANCED DATA**
- Resample the non-tumor patches to get a balanced data
- 1672 tumor images
- 1672 non-tumor images

DATA PROCESSING

- **LOOK AT THE FINAL DATA**

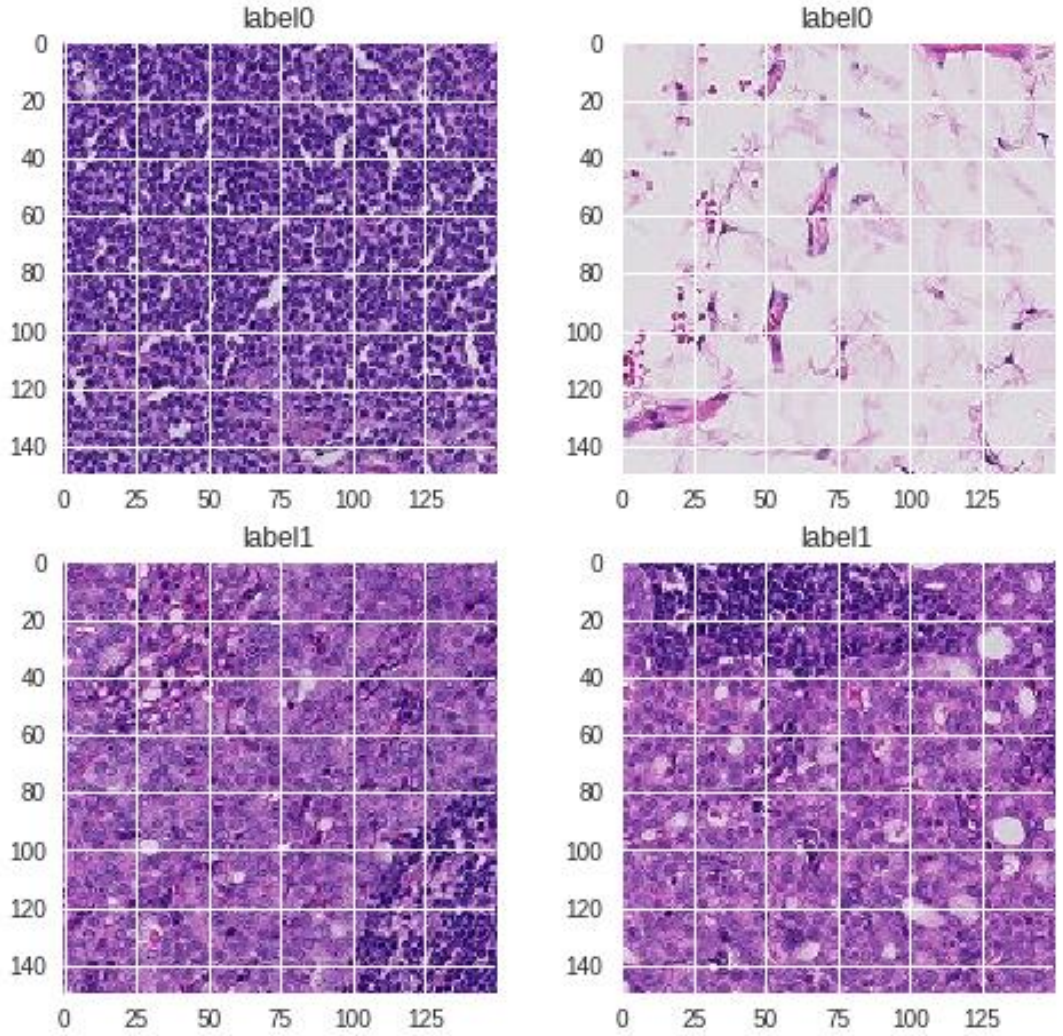

MODEL

BASE MODEL: VGG16

Add three dense layers

Not A Transfer Learning !

Because the data size is enough to train the network size of VGG.

The categories of ImageNet and cell pictures are totally different.

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 1024)	8389632
dense_2 (Dense)	(None, 512)	524800
dense_3 (Dense)	(None, 1)	513

Total params: 23,629,633
Trainable params: 23,629,633
Non-trainable params: 0

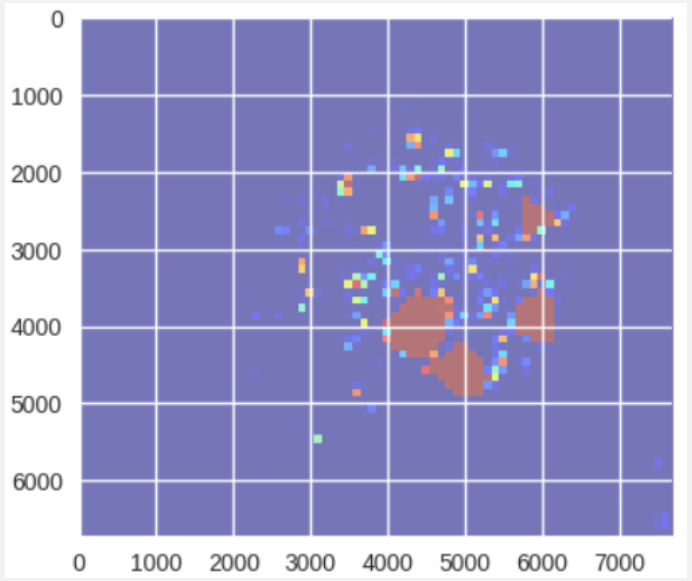

- prediction

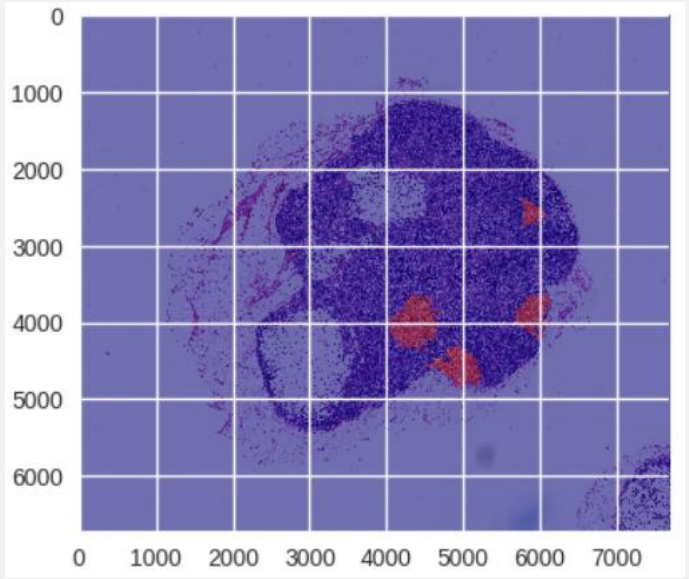

- Ground truth

EVALUATION

Heatmap Compare on the Train Image

AUC = 0.9983034004612835

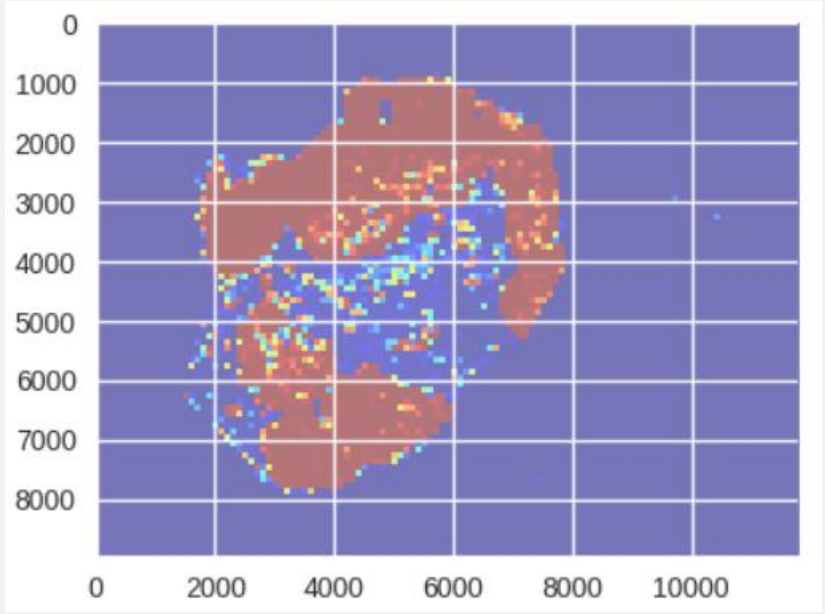

- prediction

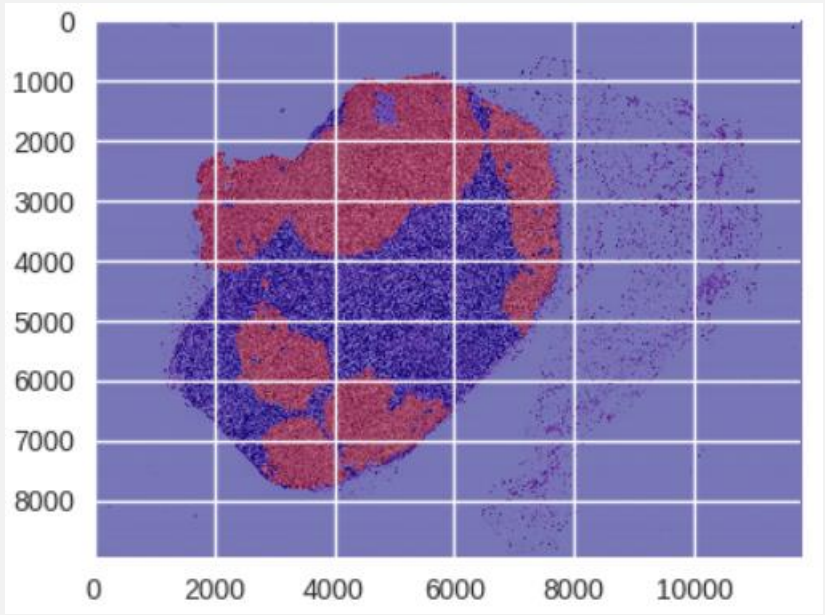

- Ground truth

EVALUATION TUMOR 110

Confusion Matrix	Ground truth NO	Ground truth YES
Prediction NO	85251678	3812081
Prediction YES	901282	15547919

Precision = 0.989538583468287

Precision = 0.957198292068494 4

F1_score = 0.9730998101842097

Test AUC = 0.97901706259915

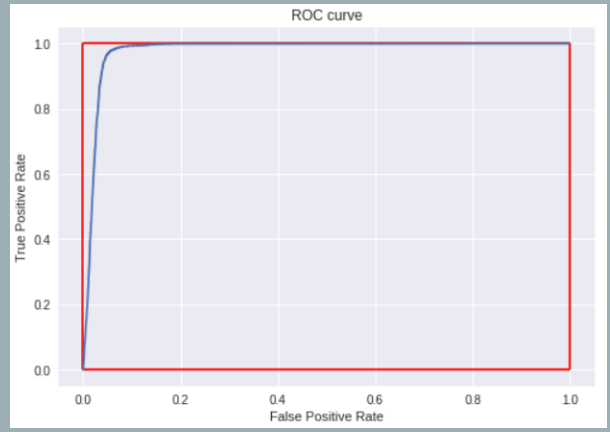

MORE DETAILS

- **ABOUT InceptionV3**
- **Because the model is so large, it's hard to train from scratch, then we must import the weights of ImageNet and freeze the model or fine tuning, but it doesn't work well.**

MORE DETAILS

- **About Optimizer and Learning Rate in my model**
- **Found the Adam didn't work well**
- **GradientDescent cannot converge well**
- **Used Momentum and tried many learning rate**

MORE DETAILS

- **Early Stop to prevent overfitting**
- **Pay attention to the training accuracy and validation accuracy**

```
Train on 2340 samples, validate on 1004 samples
Epoch 1/15
2340/2340 [==============================] - 50s 21ms/step - loss: 0.4146 - acc: 0.7893 - val_loss: 1.4903 - val_acc: 0.4910
Epoch 2/15
2340/2340 [==============================] - 40s 17ms/step - loss: 0.2859 - acc: 0.8739 - val_loss: 0.1786 - val_acc: 0.9263
Epoch 3/15
2340/2340 [==============================] - 40s 17ms/step - loss: 0.1957 - acc: 0.9188 - val_loss: 1.7294 - val_acc: 0.5329
Epoch 4/15
2340/2340 [==============================] - 40s 17ms/step - loss: 0.2661 - acc: 0.8872 - val_loss: 0.1931 - val_acc: 0.9173
Epoch 5/15
2340/2340 [==============================] - 40s 17ms/step - loss: 0.1331 - acc: 0.9513 - val_loss: 0.0894 - val_acc: 0.9671
Epoch 6/15
2340/2340 [==============================] - 40s 17ms/step - loss: 0.0665 - acc: 0.9735 - val_loss: 0.0676 - val_acc: 0.9751
Epoch 7/15
2340/2340 [==============================] - 40s 17ms/step - loss: 0.0536 - acc: 0.9812 - val_loss: 0.0338 - val_acc: 0.9890
Epoch 8/15
 512/2340 [=====>........................] - ETA: 26s - loss: 0.0086 - acc: 1.0000

KeyboardInterruptTraceback (most recent call last)
<ipython-input-32-78d7195c5599> in <module>()
```

MORE DETAILS

- **About Out of Memory**
- **I did lots of optimization in the code and release the large variable after use.**

THANK YOU!